

# Deep transfer learning with Xfer

**Andreas Damianou**  
Amazon, Cambridge UK



# Outline

---

- Deep neural networks quick reminder
- Transfer learning intro
- Xfer
  - Transfer learning via meta-learning
- Considerations

# Resources

---

- Notebook:  
[adamian.github.io/talks/Damianou\\_DL\\_Xfer.ipynb](https://adamian.github.io/talks/Damianou_DL_Xfer.ipynb)
- A more complete tutorial on deep learning:  
[adamian.github.io/talks/Damianou\\_deep\\_learning\\_rss\\_2018.pdf](https://adamian.github.io/talks/Damianou_deep_learning_rss_2018.pdf)

# Deep neural networks: hierarchical function definitions

---

A neural network is a composition of functions (layers), each parameterized with a *weight vector*  $\mathbf{w}_l$ . E.g. for 2 layers:

$$f_{\text{net}} = h_2(h_1(\mathbf{x}; \mathbf{w}_1); \mathbf{w}_2).$$

# Deep neural networks: hierarchical function definitions

---

A neural network is a composition of functions (layers), each parameterized with a *weight vector*  $\mathbf{w}_l$ . E.g. for 2 layers:

$$f_{\text{net}} = h_2(h_1(\mathbf{x}; \mathbf{w}_1); \mathbf{w}_2).$$

Generally  $f_{\text{net}} : \mathbf{x} \mapsto \mathbf{y}$  with:

$$\mathbf{h}_1 = \varphi(\mathbf{x}\mathbf{w}_1 + b_1)$$

$$\mathbf{h}_2 = \varphi(\mathbf{h}_1\mathbf{w}_2 + b_2)$$

...

$$\hat{\mathbf{y}} = \varphi(\mathbf{h}_{L-1}\mathbf{w}_L + b_L)$$

$\varphi$  is the (non-linear) activation function.

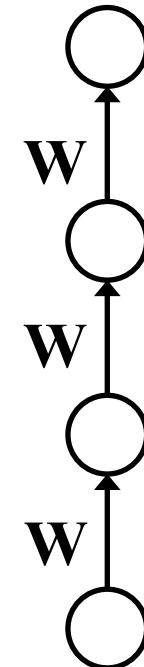
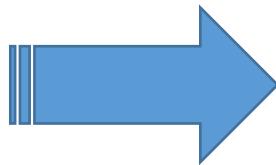
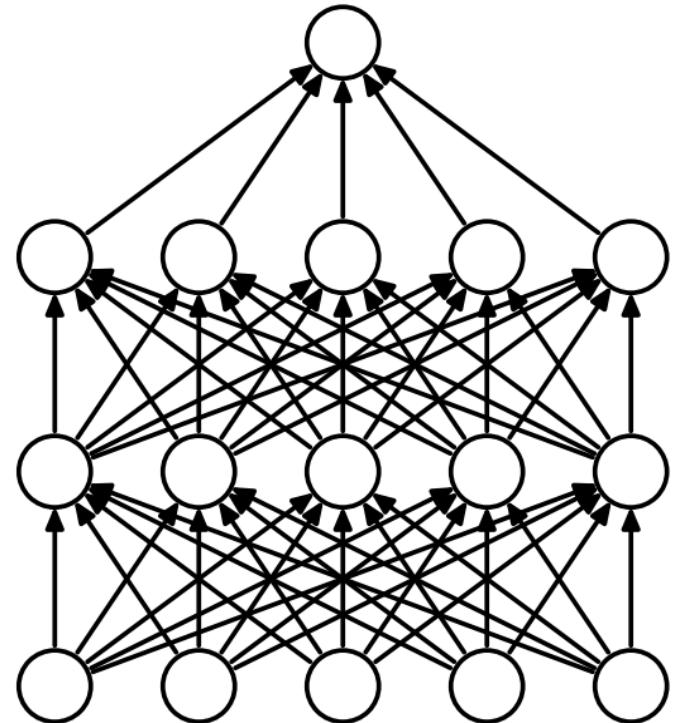
# Defining the loss

---

- We have our function approximator  $f_{\text{net}}(x) = \hat{y}$
- We have to define our loss (objective function) to relate this function outputs to the observed data.
- E.g. squared difference  $\sum_n (y_n - \hat{y}_n)^2$  or cross-entropy

# Graphical depiction

---



# Optimization and implementation

---

- Optimization done with back-propagation, based on the chain rule

*GOTO notebook!!*

# Taming the dragon

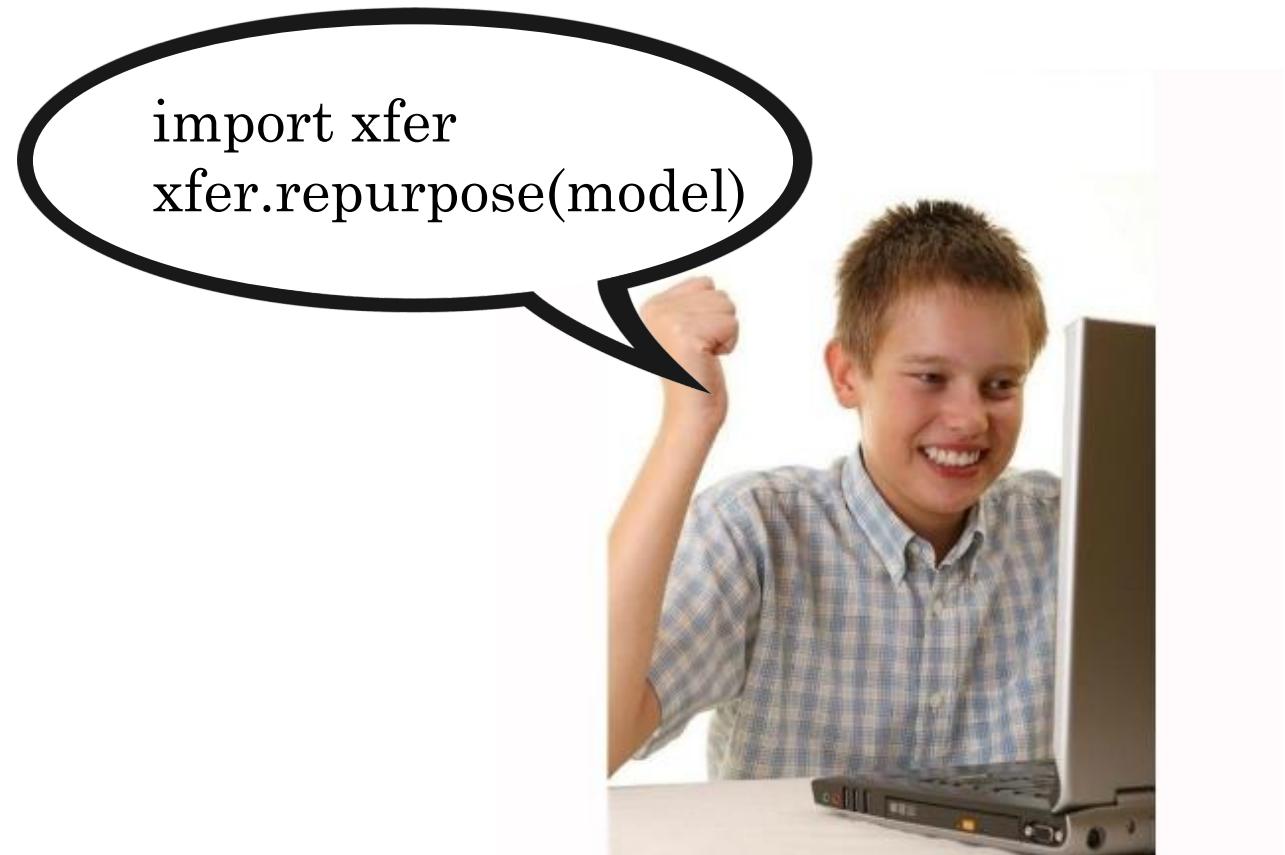
---



# Motivations for TL: DNN training requires expertise

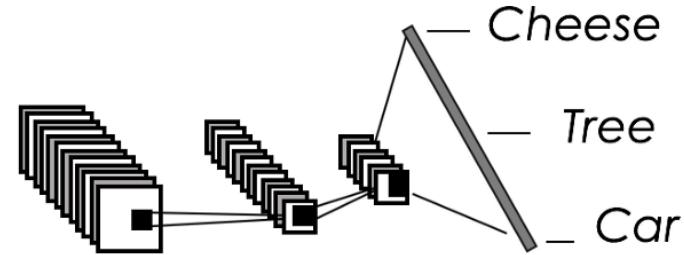
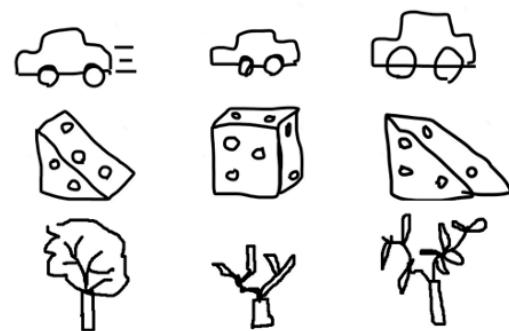
---

- Leveraging the power of DNNs even without too much expertise



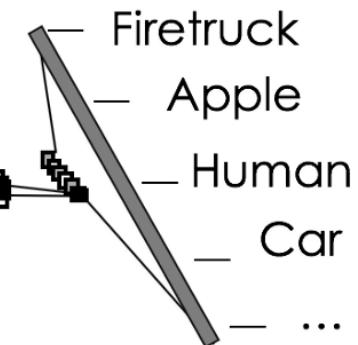
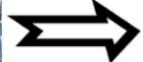
# Motivations for TL: Leverage commonalities in data

Target Task (Few images)

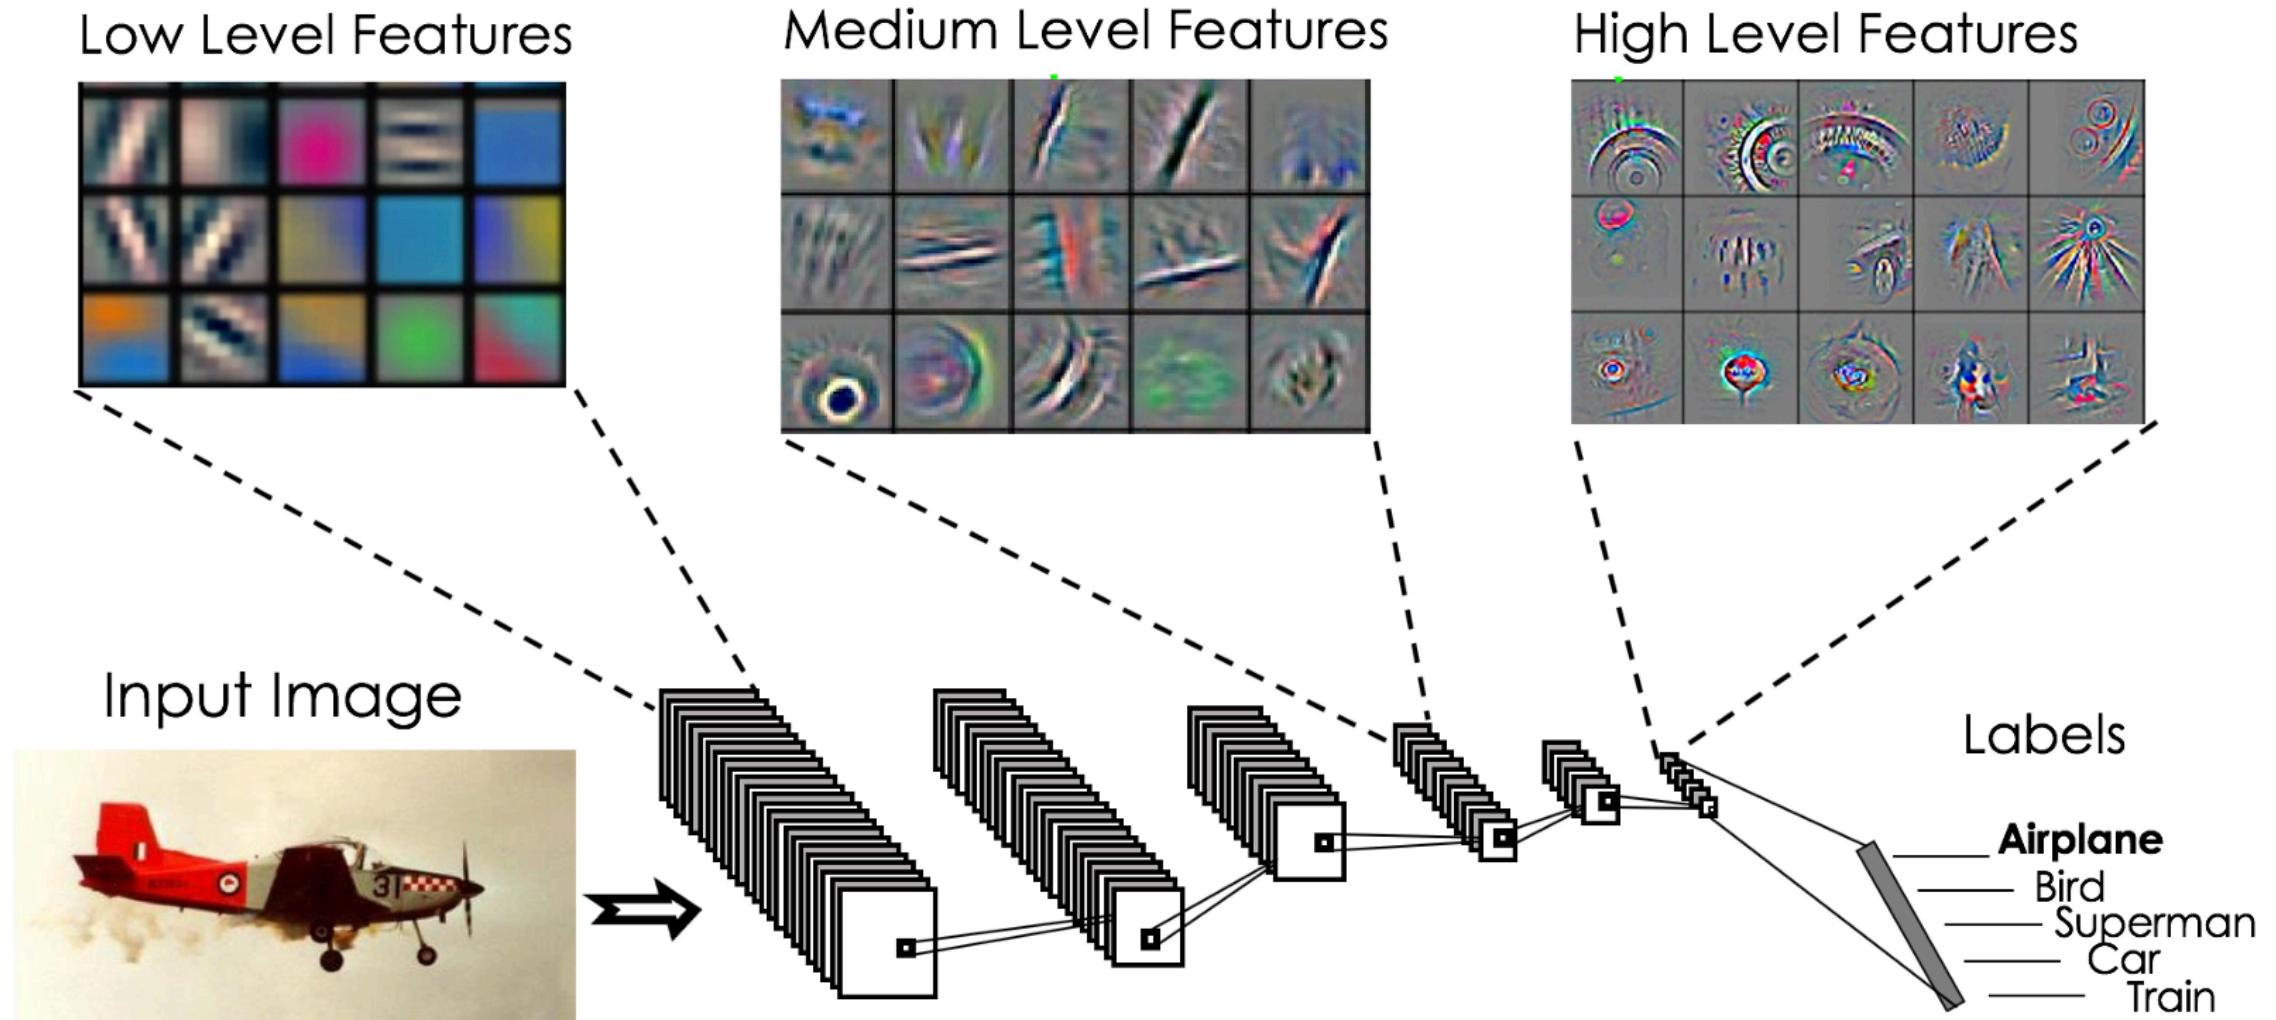


Transfer

Source Task (Many images)



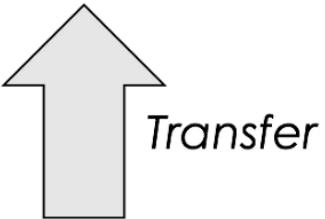
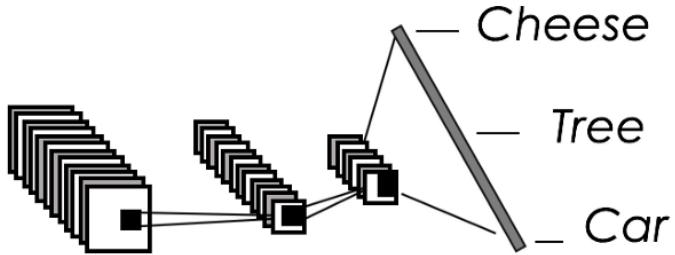
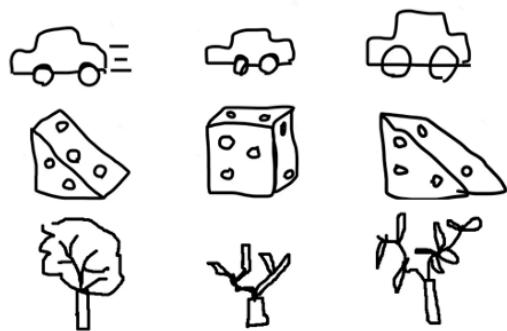
# Why does Transfer Learning work?



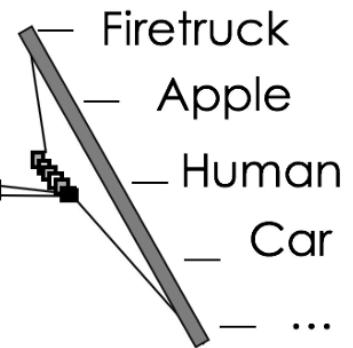
# Back to our transfer example

---

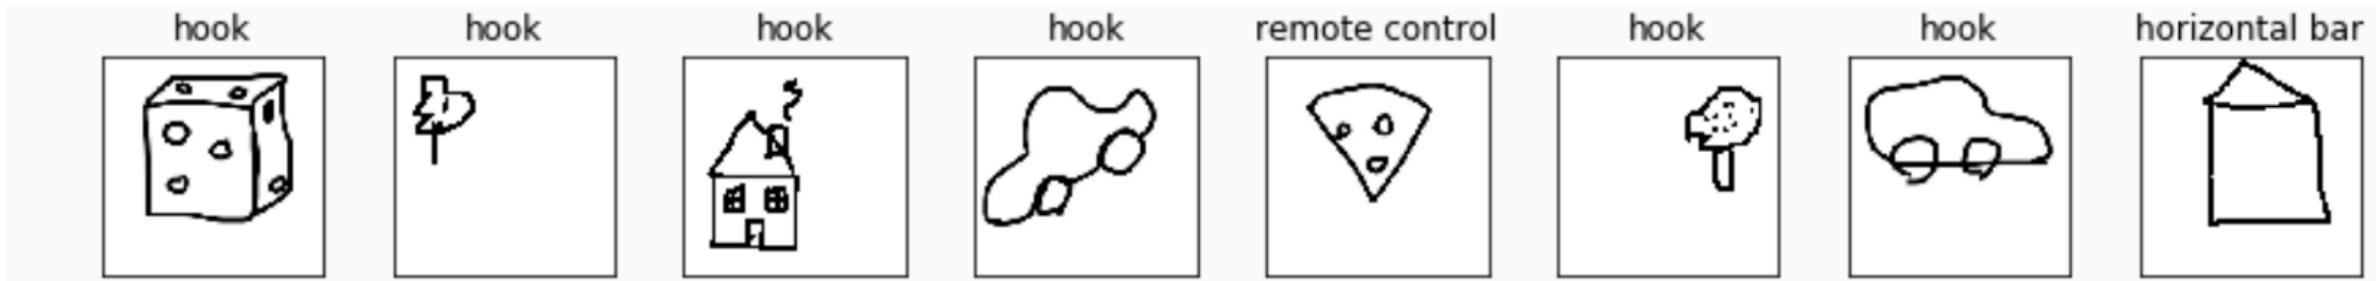
Target Task (Few images)



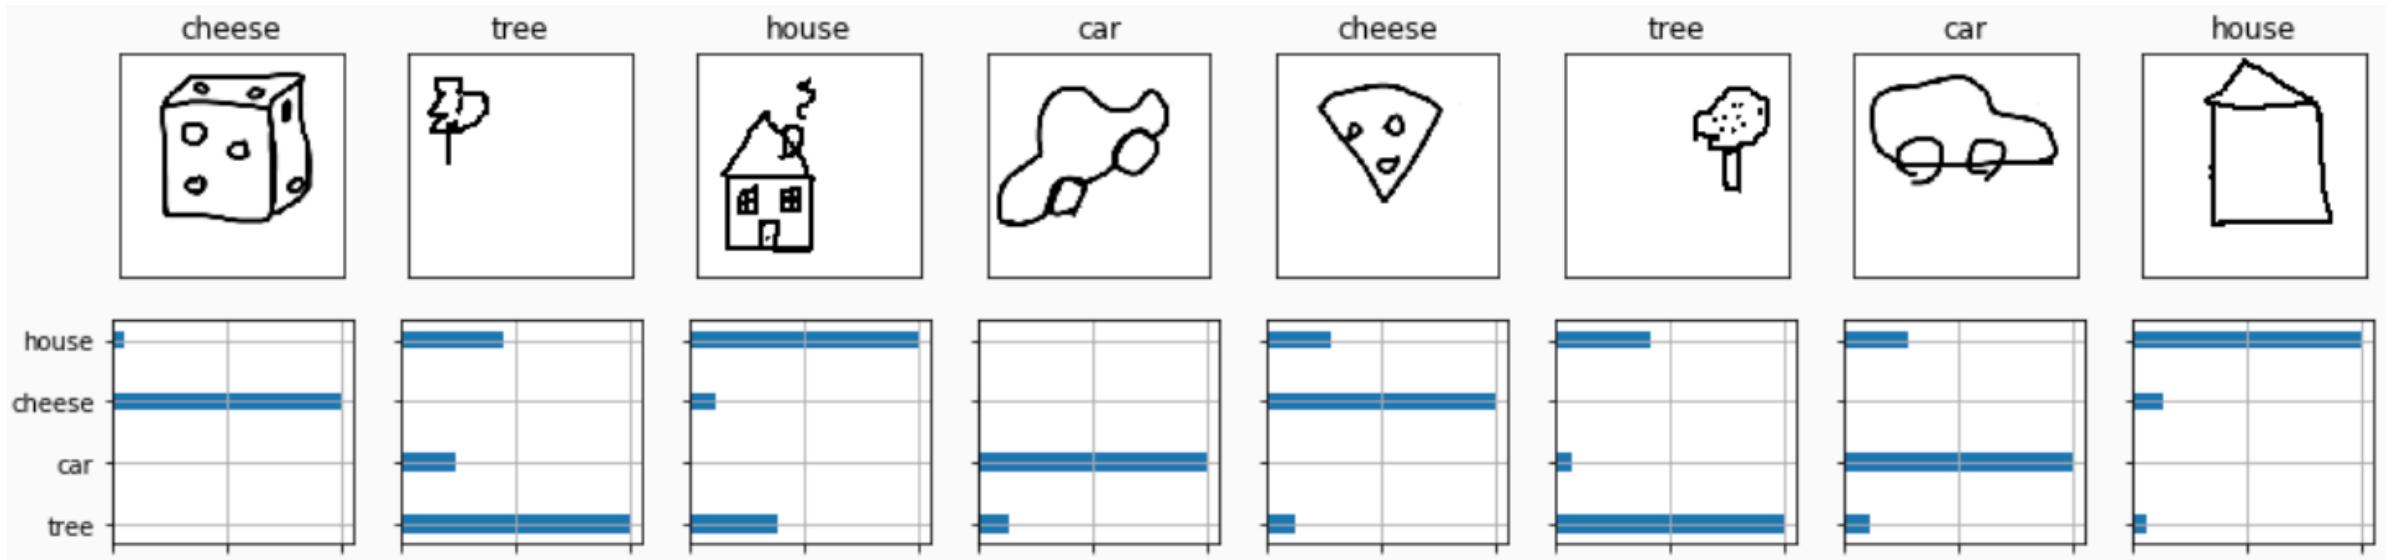
Source Task (Many images)



## Predictions using a pre-trained model (no transfer)



## Predictions using Xfer





## Deep Transfer Learning for MXNet

---

[build](#) [passing](#) [docs](#) [passing](#)  [96%](#) [pypi](#) [v1.0.0](#) [license](#) [Apache-2.0](#)

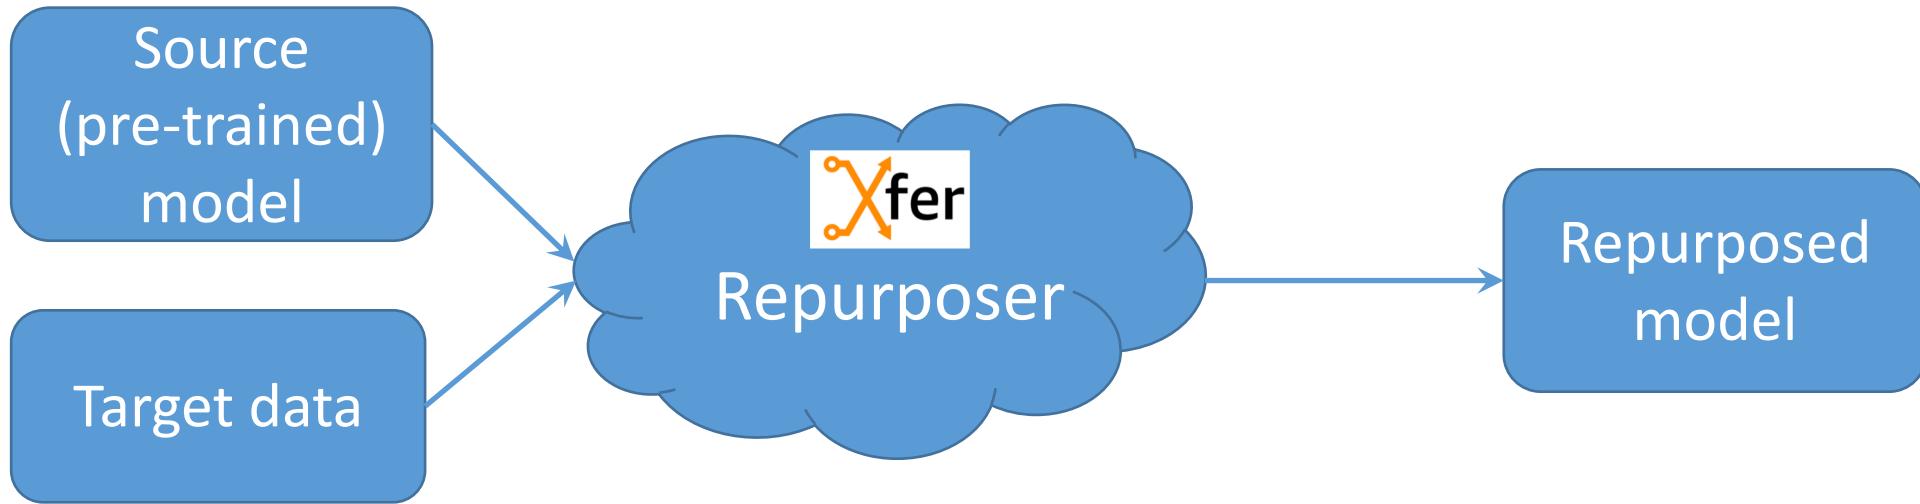
[Website](#) | [Documentation](#) | [Contribution Guide](#)

### What is Xfer?

Xfer is a library that allows quick and easy transfer of knowledge<sup>1,2,3</sup> stored in deep neural networks implemented in [MXNet](#). Xfer can be used with data of arbitrary numeric format, and can be applied to the common cases of image or text data.

# Xfer Repurposers

---



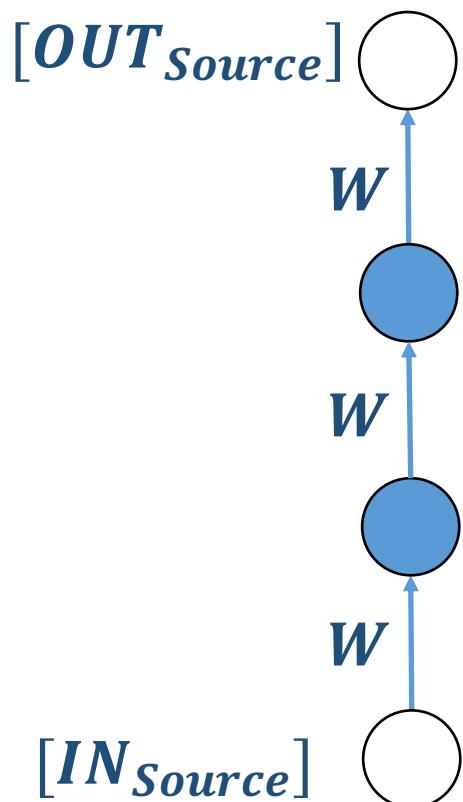
Three kinds of repurposers:

- Meta-model based
- Fine-tuning based
- Multi-task and meta-learning based (learning to learn)

# Meta-model based repurposing

---

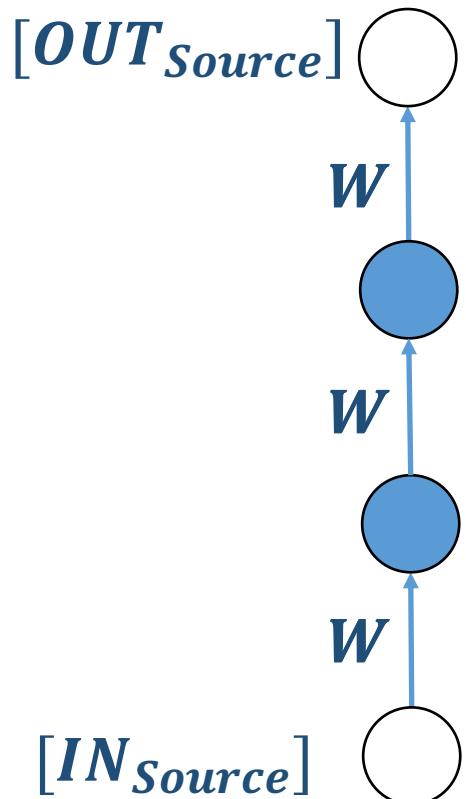
**Given:**  
**(source task)**



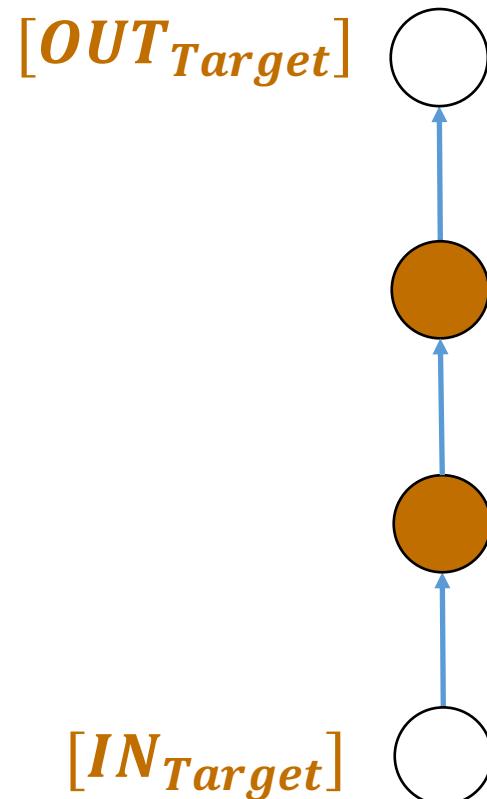
# Meta-model based repurposing

---

**Given:**  
**(source task)**



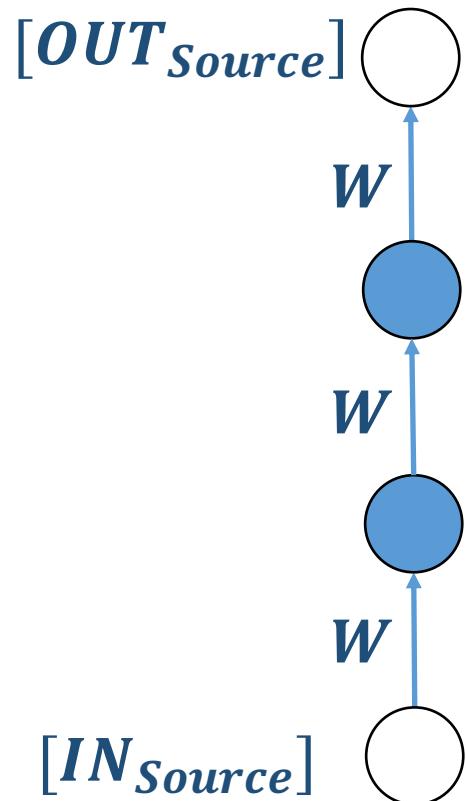
**Step 1:**  
**(target task)**



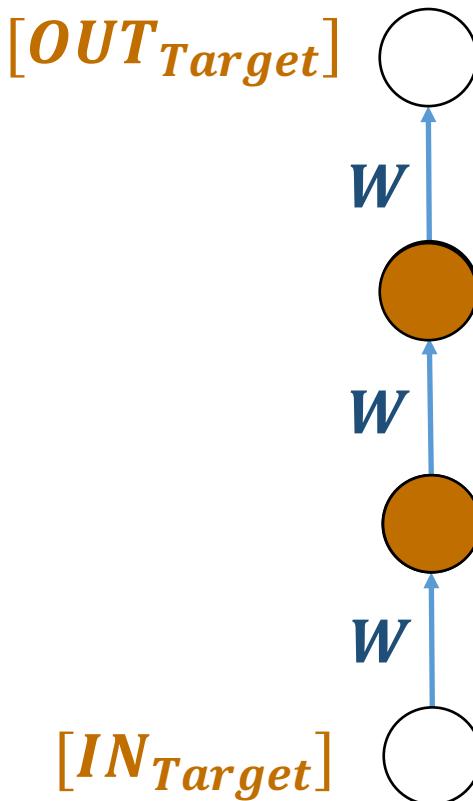
# Meta-model based repurposing

---

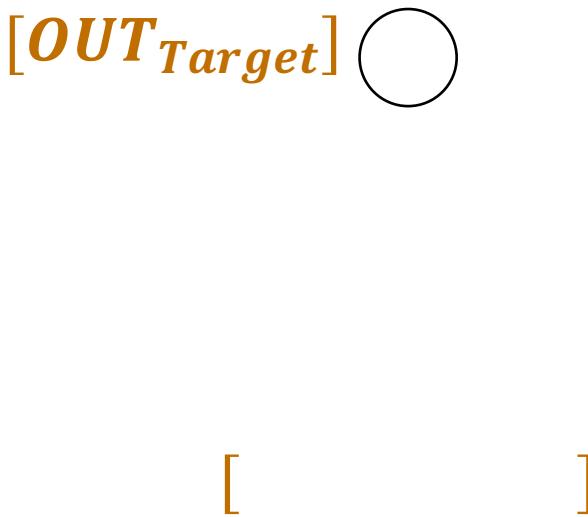
**Given:**  
**(source task)**



**Step 1:**  
**(target task)**

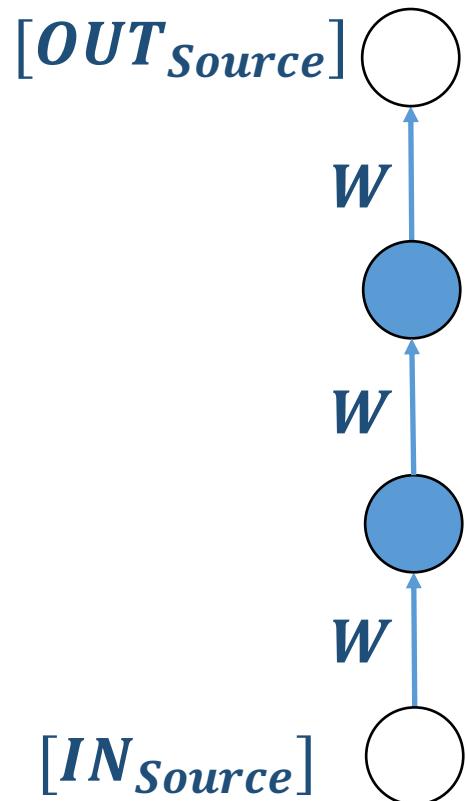


**Step 2:**  
Meta-model

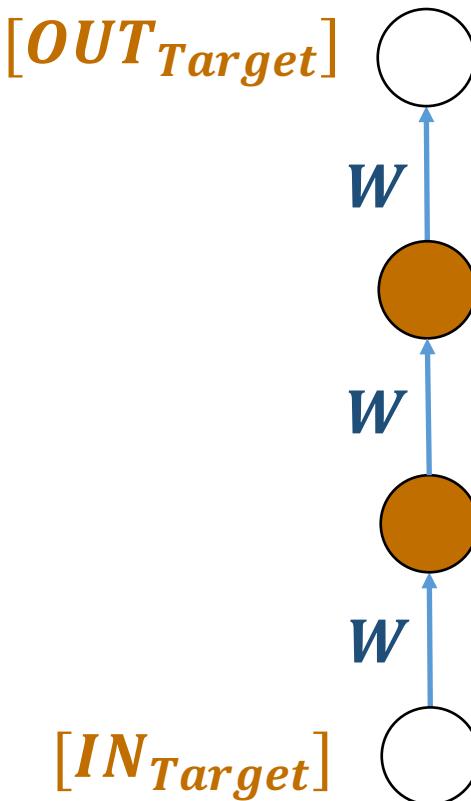


# Meta-model based repurposing

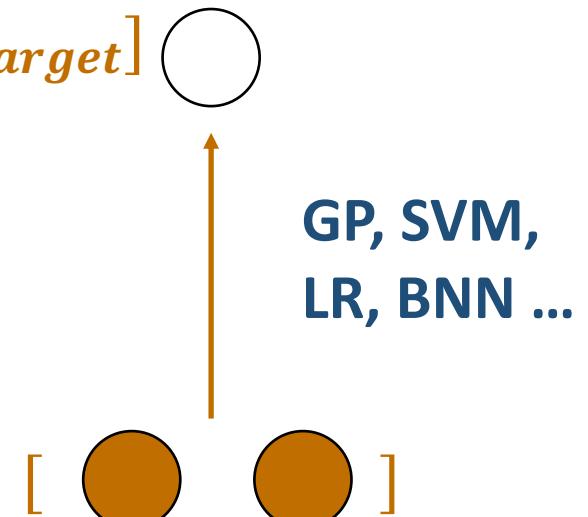
**Given:**  
**(source task)**



**Step 1:**  
**(target task)**



**Step 2:**  
Meta-model



# Meta-model based repurposing

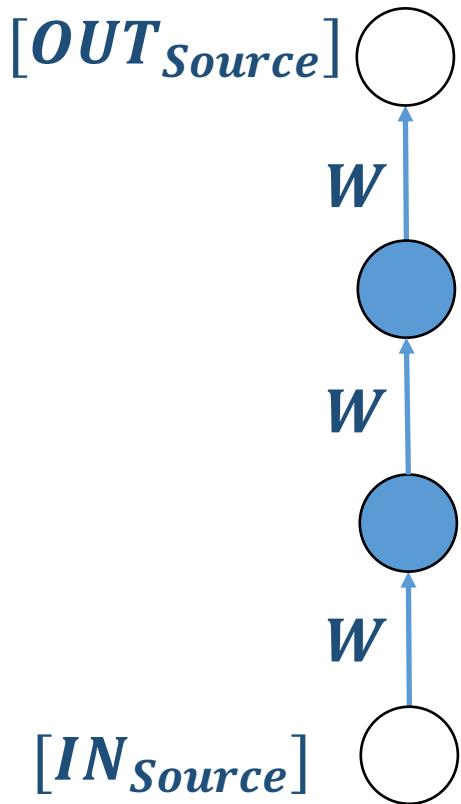
---

```
repposer = xfer.LrRepposer(source_model, feature_layer_names=['fc2','fc3'])  
  
repposer.repurpose(train_iterator)  
  
predictions = repposer.predict_label(test_iterator)
```

# Fine-tuning based repurposing

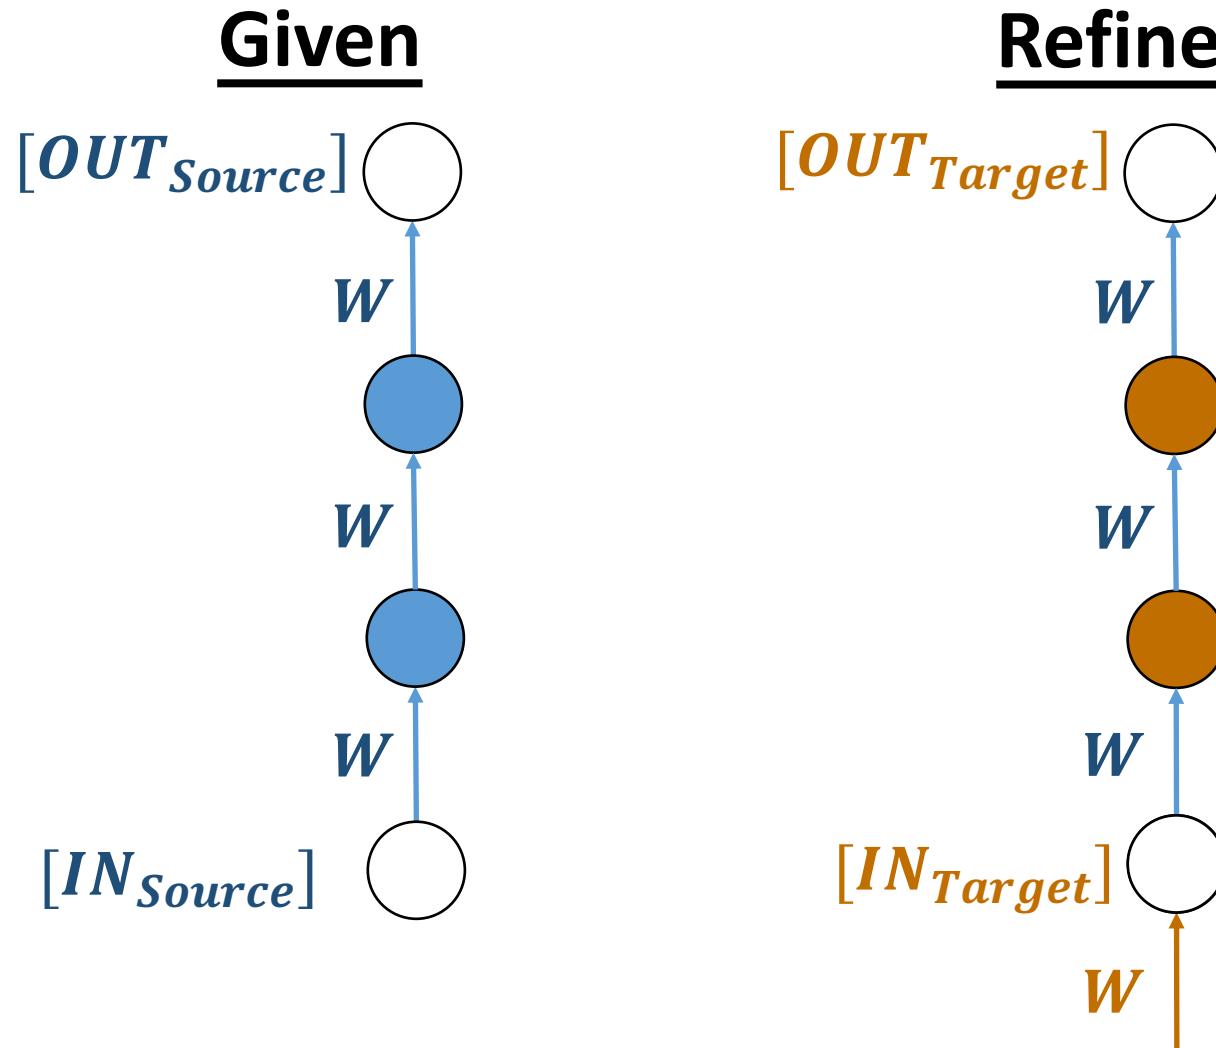
---

**Given**

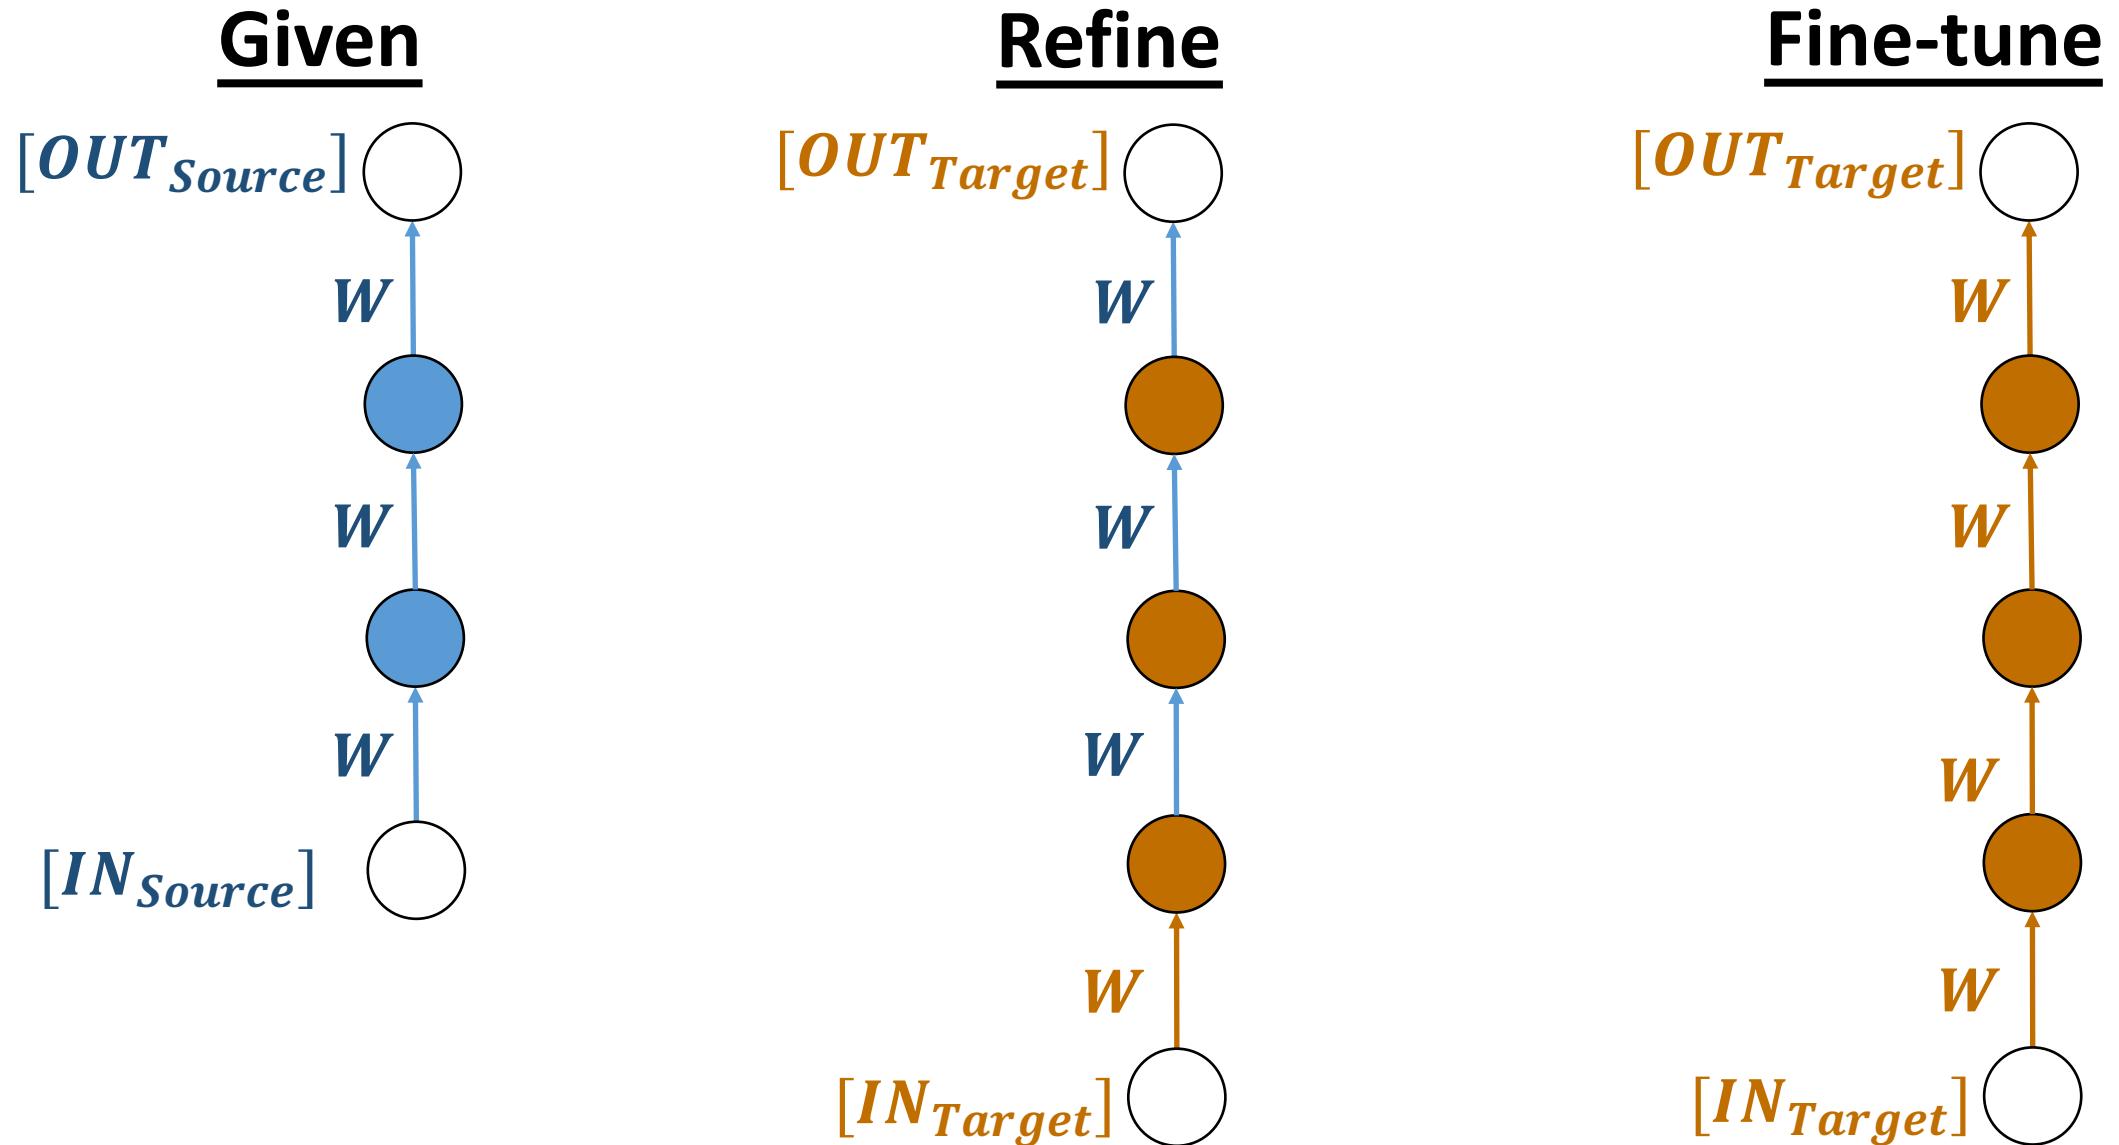


# Fine-tuning based repurposing

---



# Fine-tuning based repurposing



# Fine-tuning based repurposing

---

```
mh = xfer.model_handler.ModelHandler(source_model)

conv1 = mxnet.sym.Convolution(name='convolution1', kernel=(20,20), num_filter=64)

mh.add_layer_bottom([conv1])

mod = mh.get_module(iterator, fixed_layer_parameters=mh.get_layer_parameters(['conv1_1']),
                     random_layer_parameters=mh.get_layer_parameters(['fc6', 'fc7']))

mod.fit(iterator, num_epoch=5)
```

# Closer look at ModelHandler: inspection

---

```
mh = xfer.model_handler.ModelHandler(source_model)

print(mh.layer_names)

print(mh.get_layer_type('relu5_2'))

print(mh.get_layer_names_matching_type('Convolution'))

mh.visualize_net()
```

# Closer look at ModelHandler: feature extraction

---

```
features, labels = mh.get_layer_output(data_iterator= iterator, layer_names= ['fc6', 'fc8'])
```

# Closer look at ModelHandler: model manipulation

---

```
mh.drop_layer_top(4)

mh.drop_layer_bottom(1)

conv1 = mx.sym.Convolution(name= 'convolution1', kernel=(20,20), num_filter=64)

fc = mx.sym.FullyConnected(name= 'fullyconncted1', num_hidden= 4)

softmax = mx.sym.SoftmaxOutput(name = 'softmax')

mh.add_layer_bottom([conv1])

mh.add_layer_top([fc, softmax])
```

# Custom repurposers

```
class KNNRepurposer(xfer.MetaModelRepurposer):
    def __init__(...):
        super(KNNRepurposer, self).__init__(...)

    def _train_model_from_features(...):
        lin_model = KNeighborsClassifier(n_neighbors=self.n_neighbors,...)
        ...

    def _predict_probability_from_features(): ...

    def _predict_label_from_features(): ...

    def get_params(self): ...

    def serialize(self, file_prefix): ...
```

<https://xfer.readthedocs.io/en/master/demos/xfer-custom-repurposers.html>

# Custom repurposers

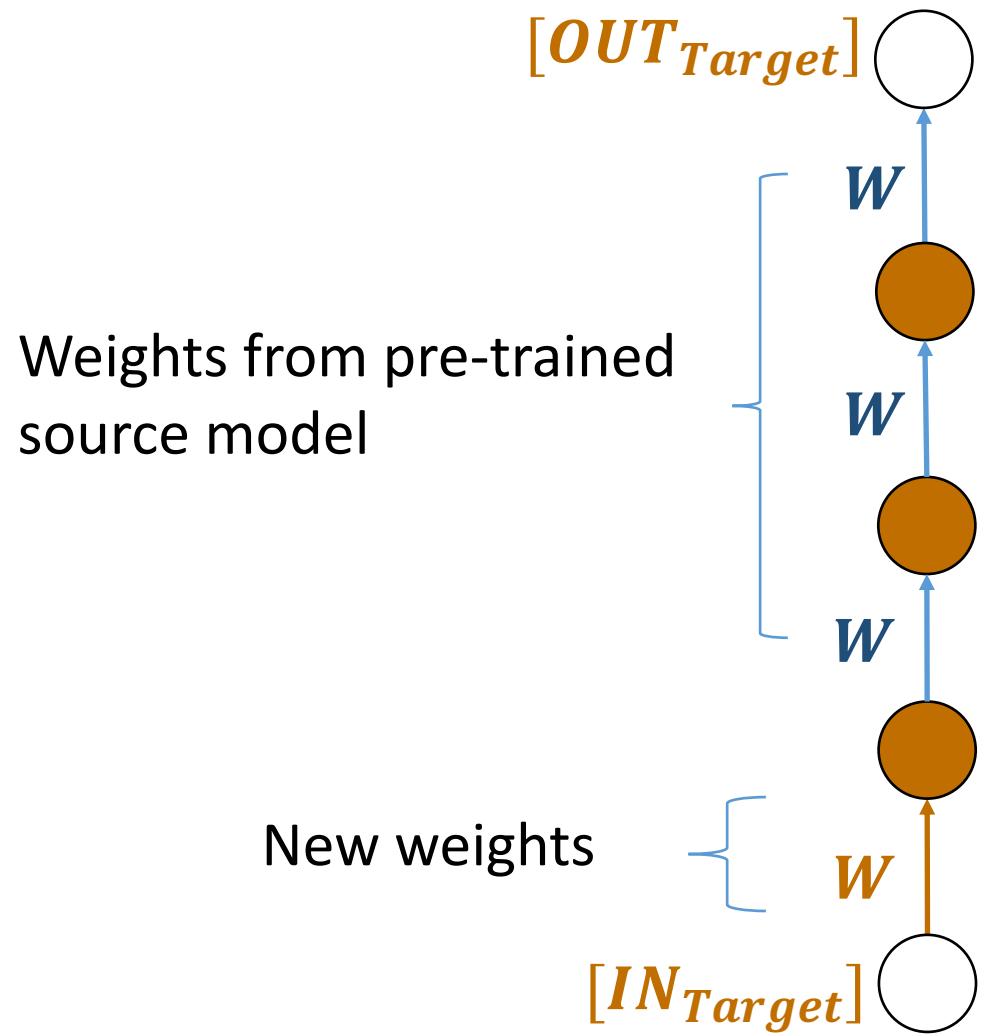
---

```
class Add2FullyConnectedRepurposer(xfer.NeuralNetworkRepurposer):  
    ...  
  
    def _create_target_module(self, train_iterator: mx.io.Dataliter):  
        model_handler = xfer.model_handler.ModelHandler(self.source_model, ...)  
  
        # ModelHandler functionality goes here...  
  
        return model_handler.get_module(train_iterator, fixed_layer_parameters=conv_layer_params)
```

<https://xfer.readthedocs.io/en/master/demos/xfer-custom-repurposers.html>

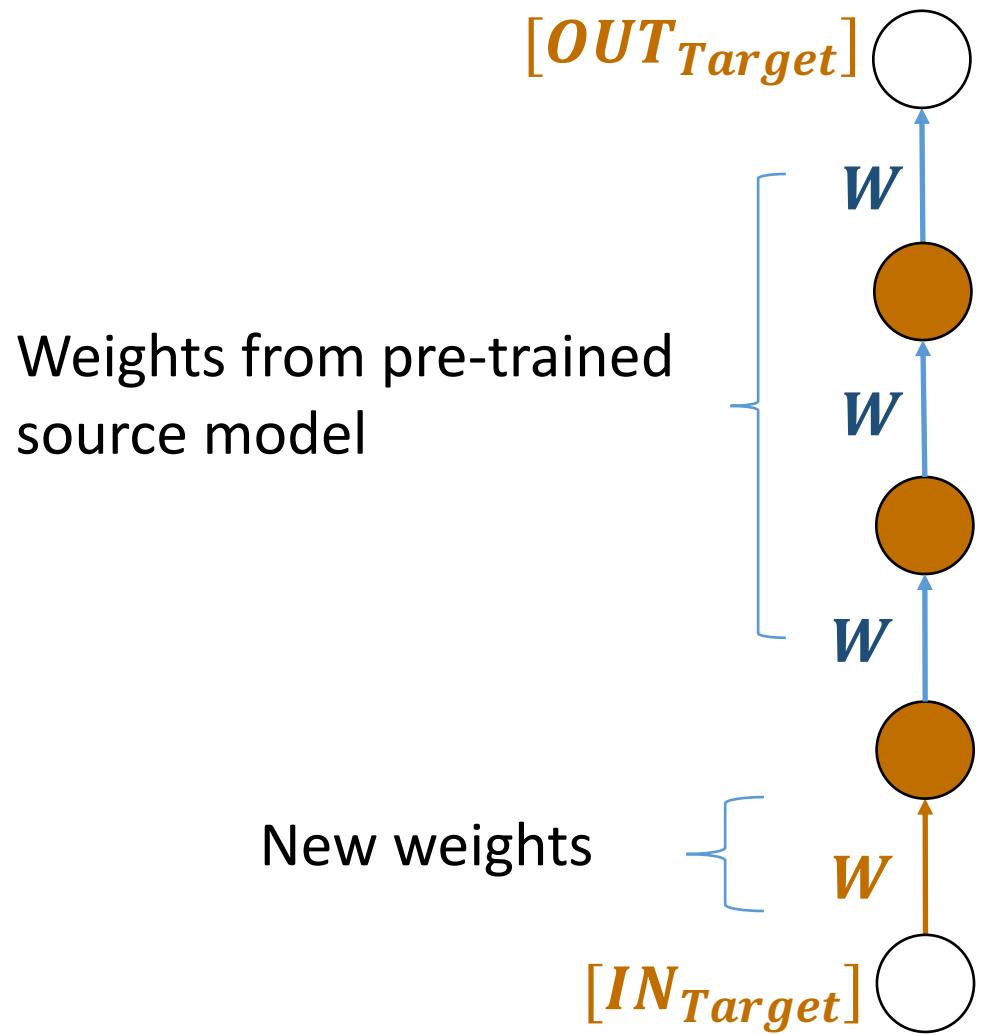
# Reminder: fine-tuning based repurposing

---



# Reminder: fine-tuning based repurposing

---



- What learning rate to use for pre-trained vs new weights?
- How many epochs?
- What optimizer to use?

# HPO for hyperparameter tuning

```
optimizer_id_to_name = {1: 'sgd', 2:'adam'}
```

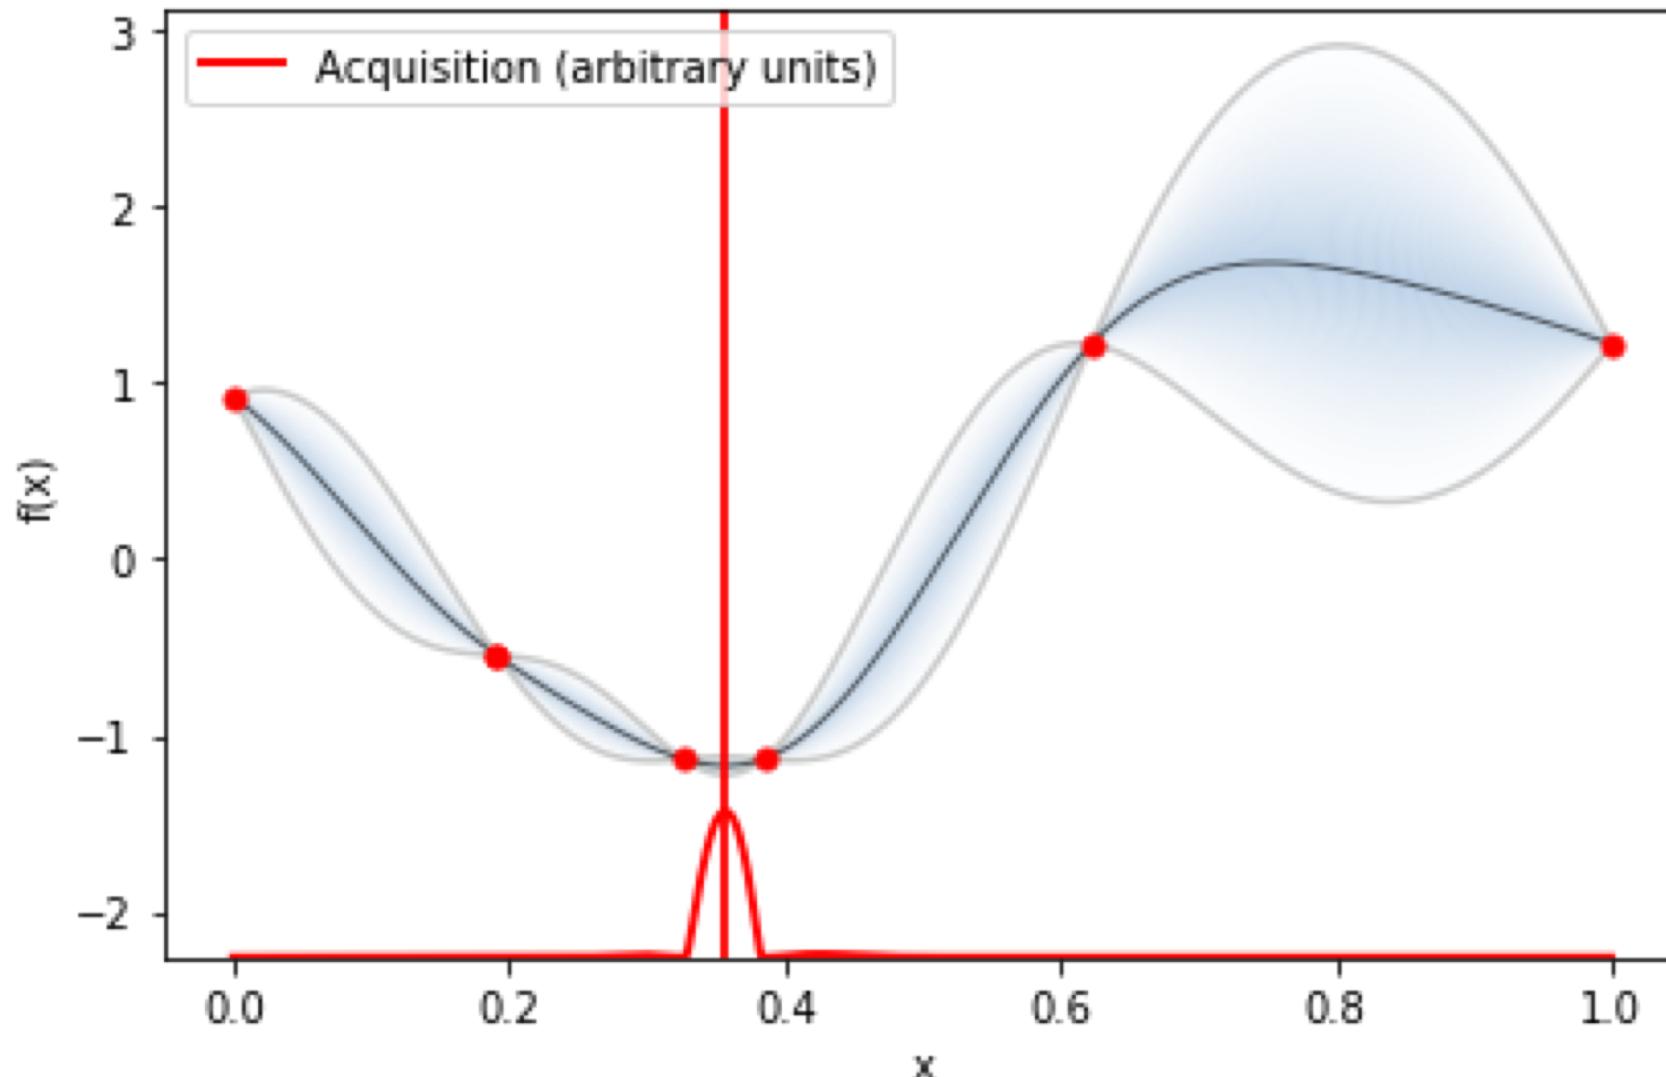
```
domain_with_2_hyperparams =  
    [ {'name': 'learning_rate', 'type': 'continuous', 'domain': (0,1)},  
      {'name': 'optimizer', 'type': 'discrete', 'domain': (1,2)}]
```

```
hyperparameter_optimizer2 = GPyOpt.methods.BayesianOptimization(  
    f = hpo_objective_function,  
    domain = domain_with_2_hyperparams))
```

```
hyperparameter_optimizer2.run_optimization()
```

<https://xfer.readthedocs.io/en/master/demos/xfer-hpo.html>

```
hyperparameter_optimizer.plot_acquisition()
```



# Xfer with Gluon

---

- Gluon models can be used with Xfer provided they use HybridBlocks so that the symbol can be extracted.

```
net = gluon.nn.HybridSequential()  
...  
net.hybridize()
```

# Xfer with Gluon

---

- Gluon models can be used with Xfer provided they use HybridBlocks so that the symbol can be extracted.

```
net = gluon.nn.HybridSequential()  
...  
net.hybridize()
```

- The Gluon model (block) is then converted into a model (symbol)

```
sym = block(data)  
args, auxs = block2symbol(block.collect_params())  
model = symbol2model(sym, data)  
model.set_params(args, auxs)
```

# Transfer through meta-learning

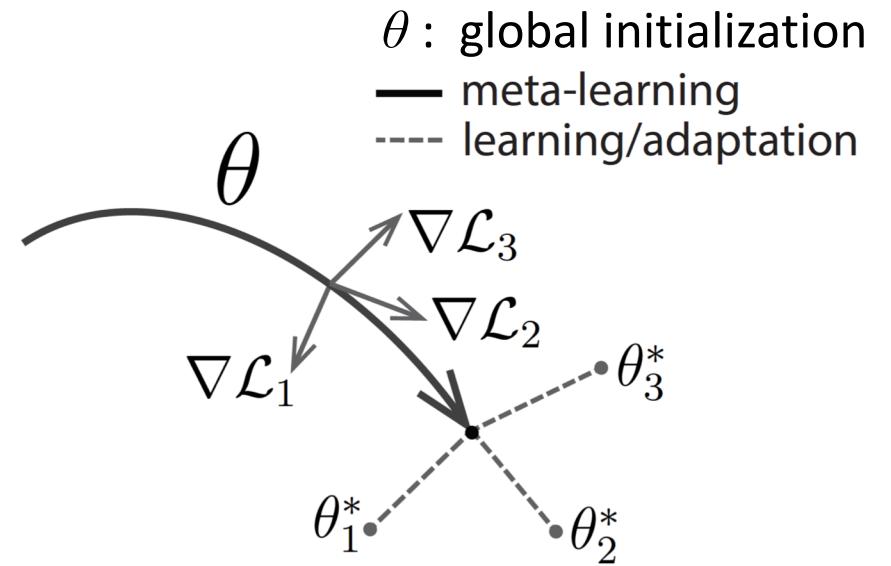
---

- Learning to learn
- Related to multi-task learning
- Our approach: transfer knowledge across learning *processes*
  - Transfer learning in a higher level of abstraction
  - Transfer learning among typically many tasks
  - All task sub-models act as source and target models

# Meta-learning or multi-task learning

---

- Optimize  $\theta$  such that on average  $\theta_i^*$  are as best as possible.

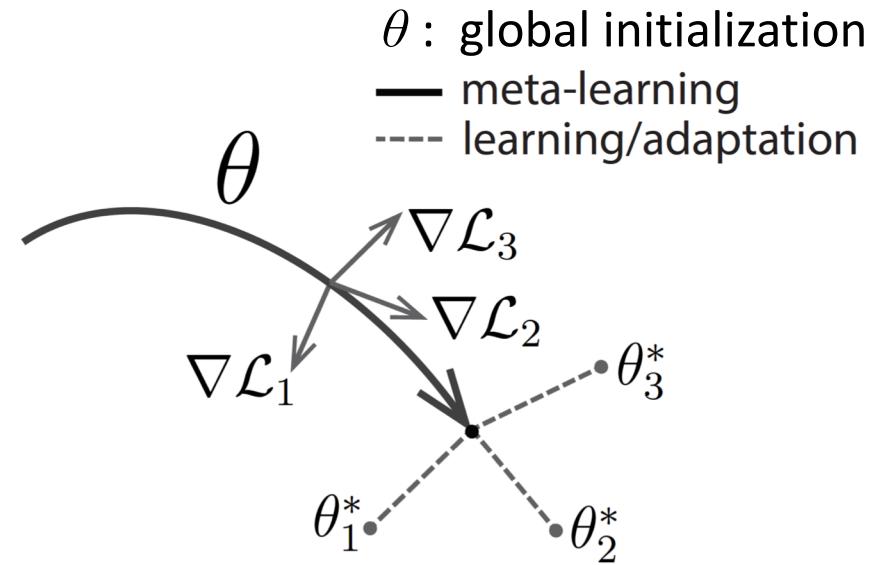


*MAML approach by Chelsea Finn et al. 2017*

# Meta-learning or multi-task learning

---

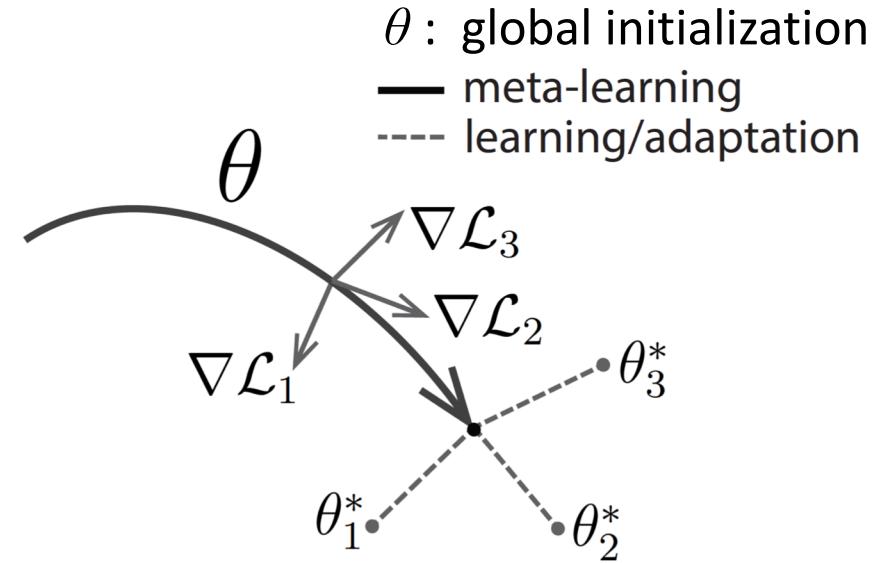
- Optimize  $\theta$  such that on average  $\theta_i^*$  are as best as possible.
- $\theta$  and  $\theta_i^*$  are in the same space.  
So we can backprop.



MAML approach by Chelsea Finn et al. 2017

# Meta-learning or multi-task learning

- Optimize  $\theta$  such that on average  $\theta_i^*$  are as best as possible.
- $\theta$  and  $\theta_i^*$  are in the same space.  
So we can backprop.



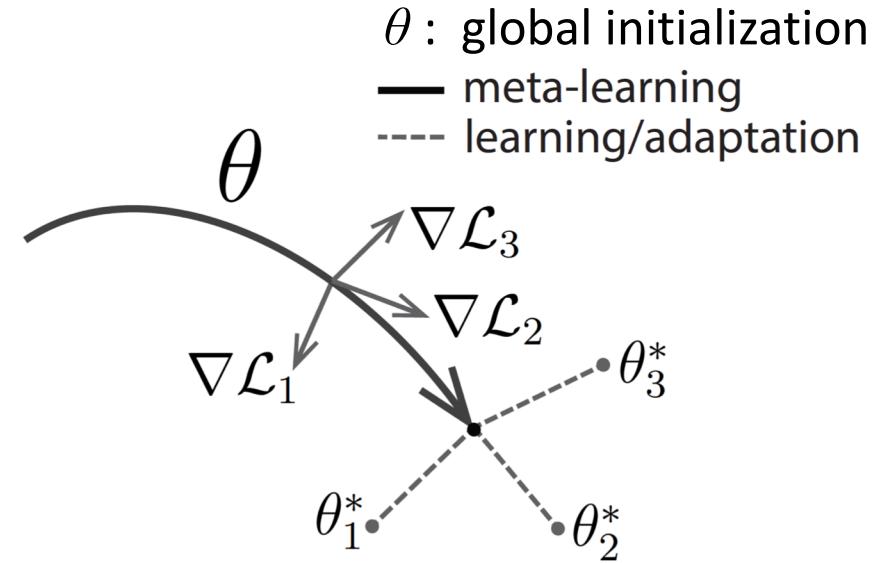
MAML approach by Chelsea Finn et al. 2017

$$\min_{\theta} \sum_{\tau_i \sim p(\tau)} \mathcal{L}_{\tau_i}(f_{\theta - \alpha \nabla_{\theta} \mathcal{L}_{\tau_i}(f_{\theta})})$$

- 
- Start with initial  $\theta$
  - for  $meta\_steps = 1, 2, \dots$  :
    - Take a batch of instances per task
    - Update  $\theta_1, \theta_2, \dots, \theta_\tau$  using each task's loss function individually
    - Update  $\theta$  such that the average of all tasks' losses is minimized

# Meta-learning or multi-task learning

- Optimize  $\theta$  such that on average  $\theta_i^*$  are as best as possible.
- $\theta$  and  $\theta_i^*$  are in the same space.  
So we can backprop.

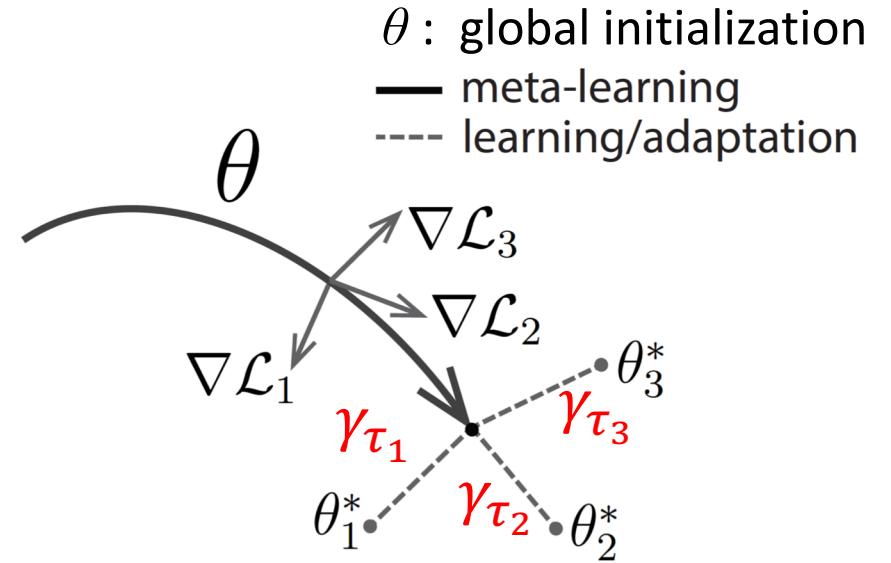


MAML approach by Chelsea Finn et al. 2017

$$\min_{\theta} \sum_{\tau_i \sim p(\tau)} \mathcal{L}_{\tau_i}(f_{\theta - \alpha \nabla_{\theta} \mathcal{L}_{\tau_i}(f_{\theta})})$$

# Meta-learning or multi-task learning

- Optimize  $\theta$  such that on average  $\theta_i^*$  is as best as possible **and**  $\theta \rightarrow \theta_i^*$  is as short as possible.
- $\theta$  and  $\theta_i^*$  are in the same space.  
So we can backprop.



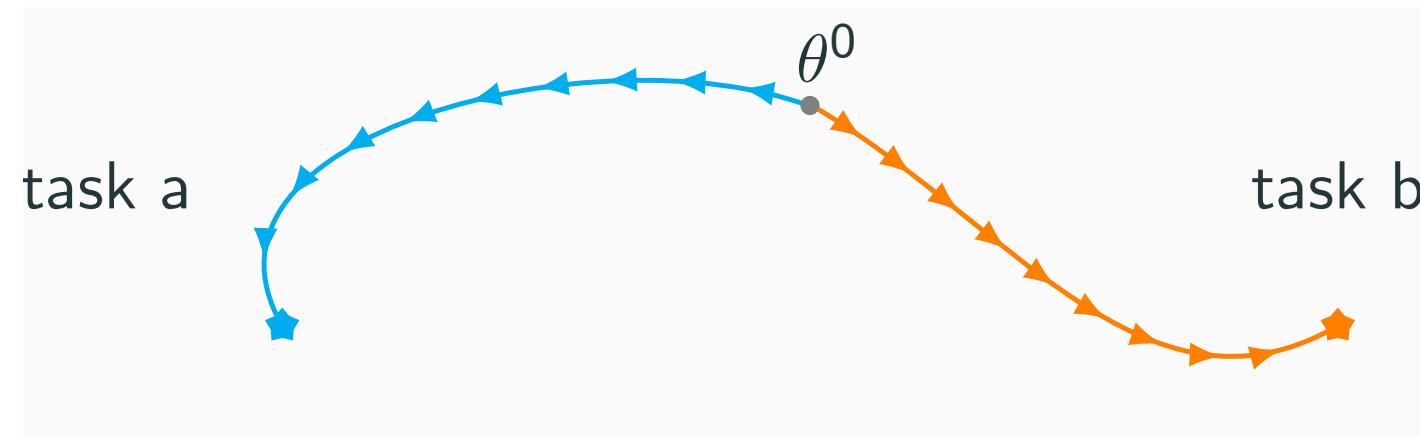
*Leap approach by Flennerhag et al. 2019  
(in Xfer soon!)*

$$\min_{\theta} \sum_{\tau_i \sim p(\tau)} \mathcal{L}_{\tau_i}(f_{\theta - \alpha \nabla_{\theta} \mathcal{L}_{\tau_i}(f_{\theta})}) + \gamma_{\tau_i}(\theta)$$

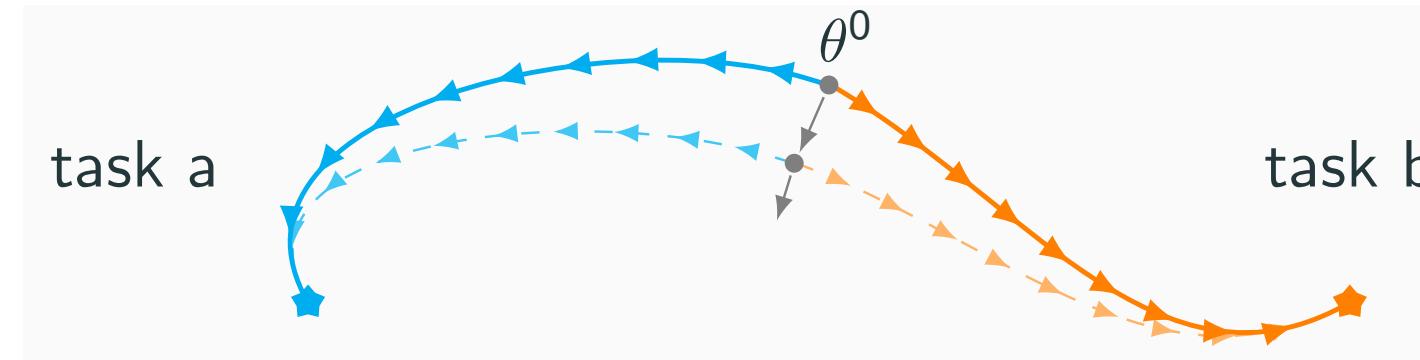
# Leap balances gradient paths from all tasks...

... to minimize the expected gradient path.

Meta-step 1



Meta-step 2



# Xfer meta-learning (*available soon!*)

```
import xfer.contrib.xfer_leap as leap

lmr = leap.leap_meta_reposer.LeapMetaRepurposer(model, num_meta_steps, num_epochs)

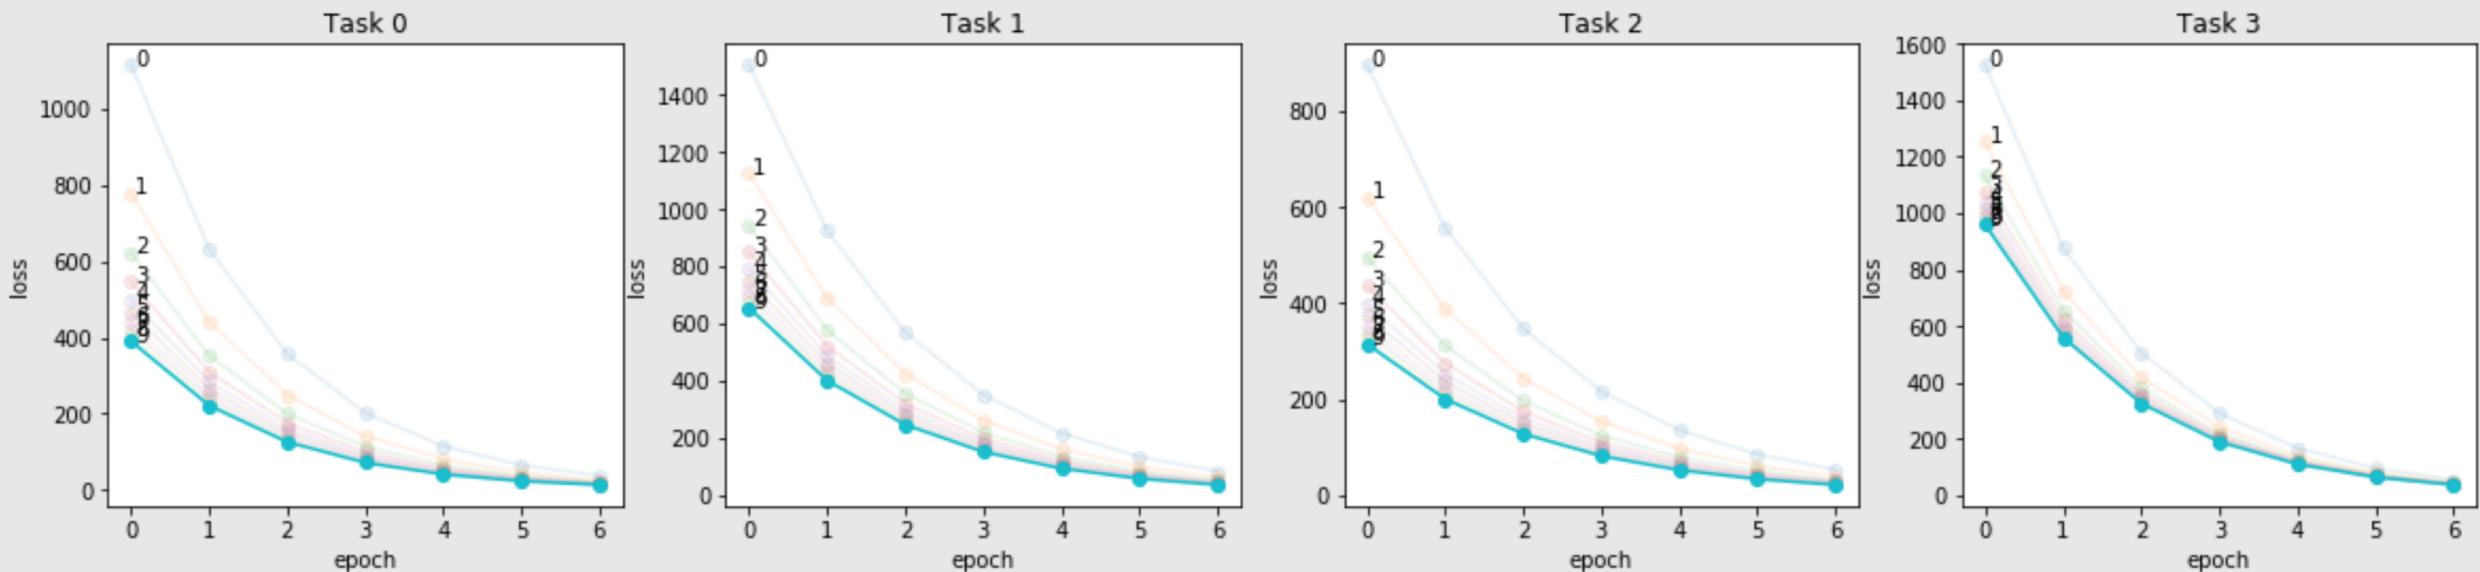
lmr.repurpose(train_data_all)
```

```
Metastep: 0, Num tasks: 4, Mean Loss: 57.061
    Metastep: 1, Task: 0, Initial Loss: 778.318, Final Loss: 25.655, Loss delta: -752.663
    Metastep: 1, Task: 1, Initial Loss: 1123.906, Final Loss: 60.993, Loss delta: -1062.913
    Metastep: 1, Task: 2, Initial Loss: 620.399, Final Loss: 38.558, Loss delta: -581.841
    Metastep: 1, Task: 3, Initial Loss: 1251.979, Final Loss: 46.972, Loss delta: -1205.006
```

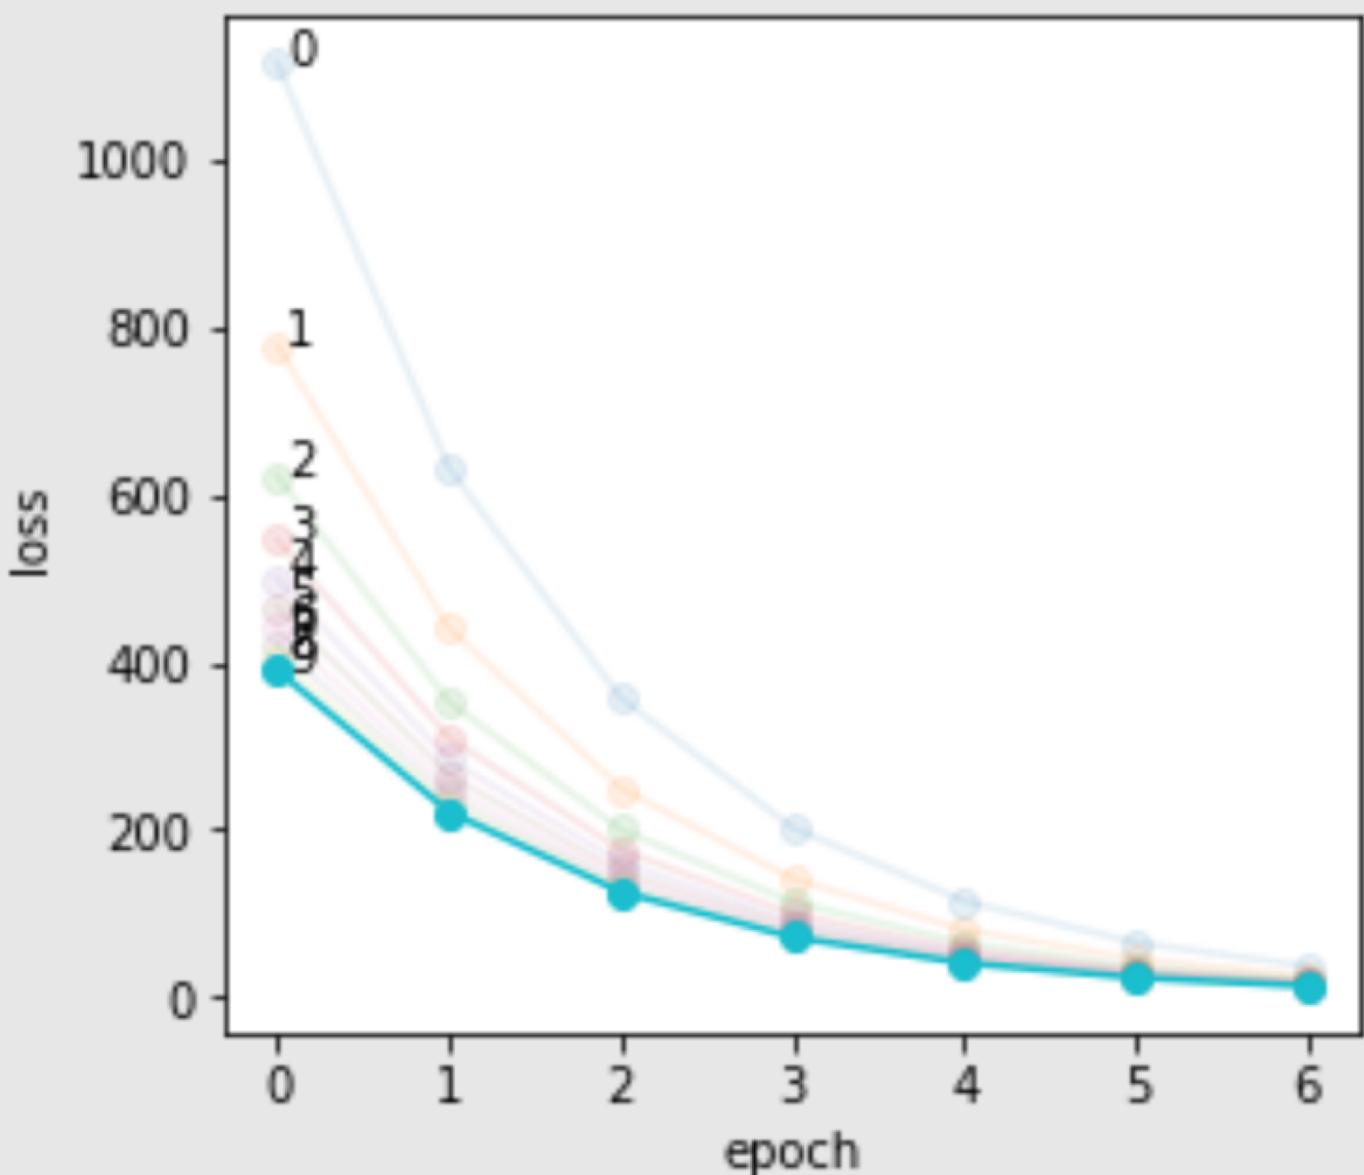
```
Metastep: 8, Num tasks: 4, Mean Loss: 27.376
    Metastep: 9, Task: 0, Initial Loss: 389.985, Final Loss: 13.036, Loss delta: -376.949
    Metastep: 9, Task: 1, Initial Loss: 654.023, Final Loss: 34.885, Loss delta: -619.138
    Metastep: 9, Task: 2, Initial Loss: 314.407, Final Loss: 21.424, Loss delta: -292.983
    Metastep: 9, Task: 3, Initial Loss: 958.127, Final Loss: 37.829, Loss delta: -920.299
```

```
lmr.meta_logger.plot_losses()
```

## Losses



## Task 0



# Data properties considerations

---

Source task:

$$\mathbf{X}_S \xrightarrow{\textit{Model}_S} \mathbf{Y}_S$$

Target task:

$$\mathbf{X}_T \xrightarrow{\textit{Model}_T} \mathbf{Y}_T$$

Transfer learning:

Use  $\textit{Model}_S$  to improve  $\textit{Model}_T$

Setting	Description	Considerations
$\mathcal{X}_S \neq \mathcal{X}_T$	Different input domains	Domain adaptation
$\mathcal{Y}_S \neq \mathcal{Y}_T$	Different label spaces	Multi-task learning might be preferable
$p(\mathbf{Y}_S) \neq p(\mathbf{Y}_T)$	Dissimilar output distribution	Transferring lower layers preferable
$p(\mathbf{X}_S) \neq p(\mathbf{X}_T)$	Dissimilar input distribution	Transferring higher layers preferable
$ \mathbf{Y}_T  \ll  \mathbf{Y}_S $	Much fewer labelled data in $T$	Data efficient TL required
$ \mathbf{Y}_T  \gg  \mathbf{Y}_S $	Much fewer labelled data in $S$	Take care of catastrophic forgetting or train T from scratch

# Acknowledgements

---

- Jordan Massiah
- Keerthana Elango
- Pablo Garcia Moreno
- Nikos Aletras
- Sebastian Flennerhag

# Thanks!

---

- Notebook:  
[adamian.github.io/talks/Damianou\\_DL\\_Xfer.ipynb](https://adamian.github.io/talks/Damianou_DL_Xfer.ipynb)
- Xfer: [github.com/amzn/xfer/](https://github.com/amzn/xfer/)
- Blog: [link.medium.com/De5BXPJ9TT](https://link.medium.com/De5BXPJ9TT)
- A more complete tutorial on deep learning:  
[adamian.github.io/talks/Damianou\\_deep\\_learning\\_rss\\_2018.pdf](https://adamian.github.io/talks/Damianou_deep_learning_rss_2018.pdf)